



SMACK Workshop Cassandra & Spark_

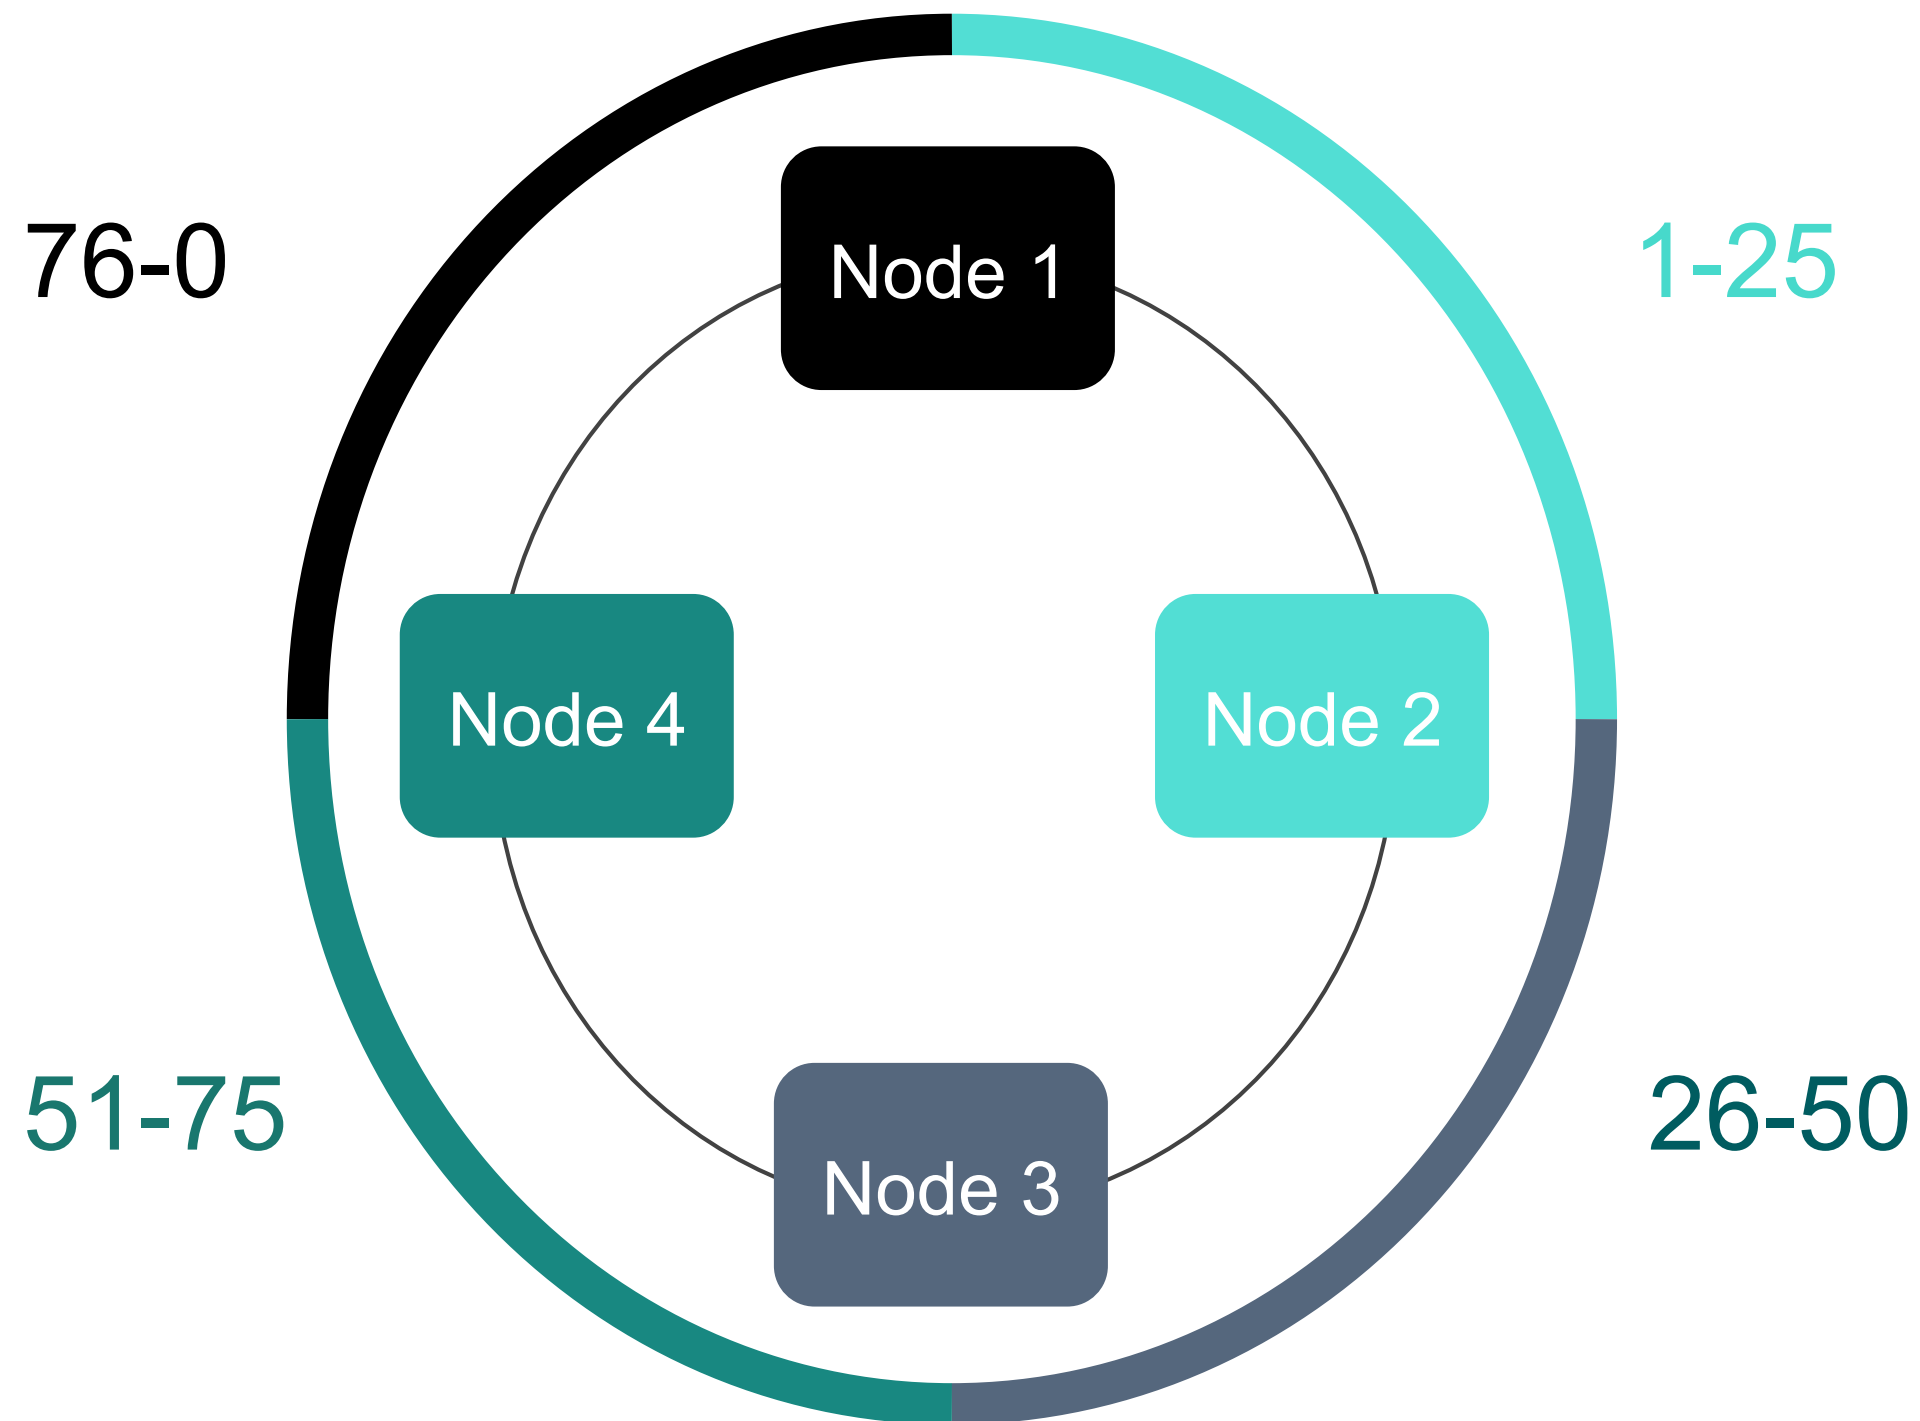
Matthias Niehoff

- **Cassandra**
- **Spark**
- **Spark & Cassandra**
- **Spark Applications**
- **Spark Streaming**
- **Spark SQL**
- **Spark MLlib**

Cassandra

- Distributed database
- Highly Available
- Linear Scalable
- Multi Datacenter Support
- No Single Point Of Failure
- CQL Query Language
 - *Similar to SQL*
 - *No Joins and aggregates*
- Eventual Consistency „Tunable Consistency“





CQL - Querying Language With Limitations_

performer
name (PK)
genre
country

SELECT * FROM performer WHERE name = 'ACDC'
→ ok

SELECT * FROM performer WHERE name = 'ACDC' and country = 'Australia'
→ not ok

SELECT country, COUNT(*) as quantity FROM artists GROUP BY country ORDER BY quantity DESC
→ not supported

- **Create tunnel to Mesos Cluster**

```
ssh -i <<pfad/zur/pem>> core@<<masterip>> -L  
9042:node-0.cassandra.mesos:9042
```

- **Start CQLSH in Virtual box**


```
/home/smack/tools/apache-cassandra-3.0.8/bin/cqlsh
```

- **Create Schema in Cassandra**

```
src/main/resources/datamodel.cql
```


Spark

- Open Source & Apache project since 2010
- Data processing Framework
 - *Batch processing*
 - *Stream processing*



- **Fast**

- *up to 100 times faster than Hadoop*
- *a lot of in-memory processing*
- *linear scalable using more nodes*

- **Easy**

- *Scala, Java and Python API*
- *Clean Code (e.g. with lambdas in Java 8)*
- *expanded API: map, reduce, filter, groupBy, sort, union, join, reduceByKey, groupByKey, sample, take, first, count*

- **Fault-Tolerant**

- *easily reproducible*

- **RDD's – Resilient Distributed Dataset**

- *Read-Only description of a collection of objects*
- *Distributed among the cluster (on memory or disk)*
- *Determined through transformations*
- *Allows automatically rebuild on failure*

- **Operations**

- *Transformations (map,filter,reduce...) → new RDD*
- *Actions (count, collect, save)*

- **Only Actions start processing!**

- **Partitions**

- *Describes the Partitions (i.e. one per Cassandra Partition)*

- **Dependencies**

- *dependencies on parent RDD's*

- **Compute**

- *The function to compute the RDD's partitions*

- **(Optional) Partitioner**

- *How is the data partitioned? (Hash, Range..)*

- **(Optional) Preferred Location**

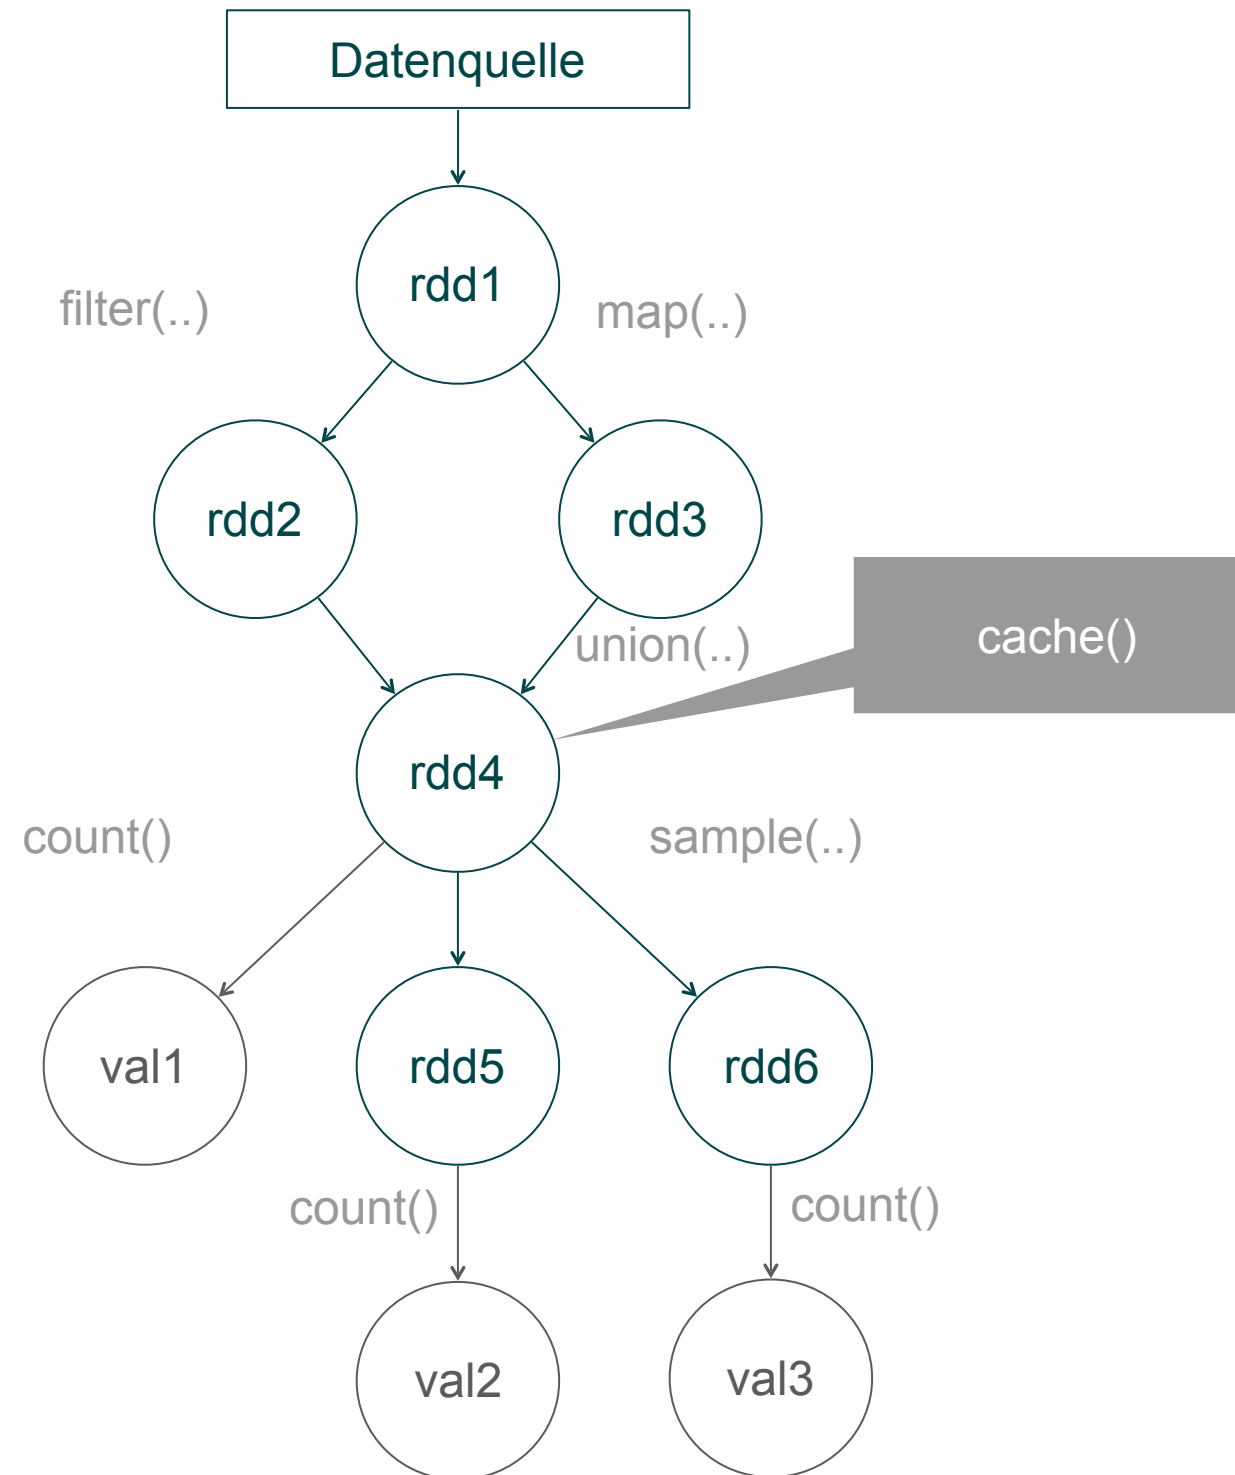
- *Where to get the data (i.e. List of Cassandra Node IP's)*

```
scala> val textFile = sc.textFile("README.md")
textFile: spark.RDD[String] = spark.MappedRDD@2ee9b6e3
```

```
scala> val linesWithSpark = textFile.filter(line =>
line.contains("Spark"))
linesWithSpark: spark.RDD[String] = spark.FilteredRDD@7dd4af09
```

```
scala> linesWithSpark.count()
res0: Long = 126
```

Reproduce RDD's Using A Tree_




- **Transformations**

- *map, flatMap*
- *sample, filter, distinct*
- *union, intersection, cartesian*

- **Actions**

- *reduce*
- *count*
- *collect, first, take*
- *saveAsTextFile*
- *foreach*

Run Spark In A Cluster_



- **DAG is built from RDD**
- **DAG Scheduler**
 - *splits graph into stages of tasks*
 - *submits each stage*
 - *resubmits failed stages*
- **TaskScheduler**
 - *launches tasks via cluster manager*
 - *resubmits failed tasks*
- **Executors**
 - *execute tasks*
 - *store and serve blocks*

Driver Node

Master Node

Worker Nodes

([atomic,collection,object] , [atomic,collection,object])

key – not unique

value

```
val fluege =  
  List( ("Thomas", "Berlin"), ("Mark", "Paris"), ("Thomas", "Madrid"))  
val pairRDD = sc.parallelize(fluege)  
  
pairRDD.filter(_._1 == "Thomas")  
  .collect  
  .foreach(t => println(t._1 + " flog nach " + t._2))
```

- **Parallelization!**

- *keys are use for partitioning*
- *pairs with different keys are distributed across the cluster*

- **Efficient processing of**


- *aggregate by key*
- *group by key*
- *sort by key*
- *joins, union based on keys*

Special OPs for PairRDDs_


`keys()`, `values()`
`mapValues(func)`, `flatMapValues(func)`
`lookup(key)`, `collectAsMap()`, `countByKey()`

`reduceByKey(func)`, `foldByKey(zeroValue)(func)`
`groupByKey()`, `cogroup(otherDataset)`
`sortByKey([ascending])`
`join(otherDataset)`, `leftOuterJoin(otherDataset)`,
`rightOuterJoin(otherDataset)`
`union(otherDataset)`, `subtractByKey(otherDataset)`


„Narrow“ (pipeline-able)



map, filter




union




join on co partitioned data

„Wide“ (shuffle)




*groupBy
on non partitioned data*




join on non co partitioned data

Spark & Cassandra


Use Spark And Cassandra In A Cluster_



DC1 - Online



DC2 - Analytics



- **Spark Cassandra Connector by Datastax**
 - *<https://github.com/datastax/spark-cassandra-connector>*
- **Cassandra tables as Spark RDD (read & write)**
- **Mapping of C* tables and rows onto Java/Scala objects**
- **Server-Side filtering („where“)**
- **Compatible with**
 - *Spark ≥ 0.9*
 - *Cassandra ≥ 2.0*
- **Clone & Compile with SBT or download at Maven Central**

- **Read complete table**

```
val movies = sc.cassandraTable("movie", "movies")  
// returns CassandraRDD[CassandraRow]
```

- **Read selected columns**

```
val movies = sc.cassandraTable("movie", "movies").select("title", "year")
```

- **Filter rows**

```
val movies = sc.cassandraTable("movie", "movies").where("title = 'Die  
Hard'")
```

- **Access Columns in Result Set**

```
movies.collect.foreach(r => println(r.get[String]("title")))
```

Read As Tuple

```
val movies =  
sc.cassandraTable[(String,Int)]("movie","movies")  
.select("title","year")
```

```
val movies =  
sc.cassandraTable("movie","movies")  
.select("title","year")  
.as((_: String, _:Int))
```

// both result in a CassandraRDD[(String,Int)]

Read As Case Class

```
case class Movie(title: String, year: Int)

sc.cassandraTable[Movie]("movie", "movies").select("title", "year")

sc.cassandraTable("movie", "movies").select("title", "year").as(Movie)
```

- Every RDD can be saved

- *Using Tuples*

```
val tuples = sc.parallelize(Seq(("Hobbit", 2012), ("96 Hours", 2008)))  
tuples.saveToCassandra("movie", "movies", SomeColumns("title", "year"))
```

- *Using Case Classes*

```
case class Movie (title:String, year: int)  
val objects =  
    sc.parallelize(Seq(Movie("Hobbit", 2012), Movie("96 Hours", 2008)))  
objects.saveToCassandra("movie", "movies")
```


Pair RDDs With Cassandra_

```
// Load and format as Pair RDD
val pairRDD = sc.cassandraTable("movie", "director")
    .map(r => (r.getString("country"), r))

// Directors / Country, sorted
pairRDD.mapValues(v => 1).reduceByKey(_+_).
    .sortBy(-_._2).collect.foreach(println)

// or, unsorted
pairRDD.countByKey().foreach(println)

// ALL Countries
pairRDD.keys()
```

director		
<i>name</i>	<i>text</i>	<i>K</i>
<i>country</i>	<i>text</i>	

- Automatically on read
- Not automatically on write
 - *No Shuffling Spark Operations -> Writes are local*
 - *Shuffling Spark Operations*
 - *Fan Out writes to Cassandra*
 - *repartitionByCassandraReplica("keyspace", "table") before write*
- Joins with data locality

```
sc.cassandraTable[CassandraRow](KEYSPACE, A)
  .repartitionByCassandraReplica(KEYSPACE, B)
  .joinWithCassandraTable[CassandraRow](KEYSPACE, B)
  .on(SomeColumns("id"))
```

Spark Streaming

Stream Processing With Spark Streaming_

- Real Time Processing using micro batches
- Supported sources: TCP, S3, Kafka, Twitter,..
- Data as Discretized Stream (DStream)
- Same programming model as for batches
- All Operations of the GenericRDD & SQL & MLlib
- Stateful Operations & Sliding Windows



Spark Streaming - Example_

```
import org.apache.spark.streaming._

val ssc = new StreamingContext(sc, Seconds(1))

val stream = ssc.socketTextStream("127.0.0.1", 9999)

stream.map(x => 1).reduce(_ + _).print()

ssc.start()

// await manual termination or error
ssc.awaitTermination()

// manual termination
ssc.stop()
```


Create an Application

- Normal Scala Application
- SBT as build tool
- source in *src/main/scala-2.10*
- *assembly.sbt* in root and project directory
- *build.sbt* in root directory

```
libraryDependencies += "com.datastax.spark" % "spark-cassandra-connector" % "1.3.0"  
libraryDependencies += "org.apache.spark" % "spark-core" % "1.3.1" % "provided"  
libraryDependencies += "org.apache.spark" % "spark-mllib_2.10" % "1.3.1" % "provided"  
libraryDependencies += "org.apache.spark" % "spark-streaming_2.10" % "1.3.1" %  
"provided"
```

- *sbt assembly* to build

- **`dcos spark run --submit-args="--class org.apache.spark.examples.SparkPi https://downloads.mesosphere.com.s3.amazonaws.com/assets/spark/spark-examples_2.10-1.5.0.jar"`**


Spark Exercises

- **Streaming Application 1**
 - *Simple Count & Aggregate*
 - *Save to Cassandra*
- **Streaming Application 2**
 - *State & Window Processing*
 - *Advanced*
- **Batch Application**
 - *Read from Cassandra*
 - *Filter Data*


Spark MLlib

- **Fully integrated in Spark**
 - *Scalable*
 - *Scala, Java & Python APIs*
 - *Use with Spark Streaming & Spark SQL*
- **Packages various algorithms for machine learning**
- **Includes**
 - *Clustering*
 - *Classification*
 - *Prediction*
 - *Collaborative Filtering*
- **Still under development**
 - *performance, algorithms*

set of data points



meaningful clusters




MLLib Example - Clustering (using KMeans)_

```
// Load and parse data
val data = sc.textFile("data/mllib/kmeans_data.txt")
val parsedData = data
  .map(s => Vectors.dense(s.split(' '))
    .map(_toDouble)).cache()


// Cluster the data into 3 classes using KMeans with 20 iterations
val clusters = KMeans.train(parsedData, 2, 20)

// Evaluate clustering by computing Sum of Squared Errors
val SSE = clusters.computeCost(parsedData)
println("Sum of Squared Errors = " + WSSSE)
```


MLLib Example - Classification_



Binary classification:



Multi-class classification:



MLLib Example - Classification (Linear SVM)_

```
// Load training data in LIBSVM format.
val data =
  MLUtils.loadLibSVMFile(sc, "sample_libsvm_data.txt")


























// Split data into training (60%) and test (40%).
val splits = data.randomSplit(Array(0.6, 0.4), seed = 11L)
val training = splits(0).cache()
val test = splits(1)

// Run training algorithm to build the model
val numIterations = 100
val model = SVMWithSGD.train(training, numIterations)
```

MLLib Example - Classification (Linear SVM)_

```
// Compute raw scores on the test set.  
val scoreAndLabels = test.map { point =>  
    val score = model.predict(point.features)  
    (score, point.label)  
}  
  
// Get evaluation metrics.  
val metrics = new  
BinaryClassificationMetrics(scoreAndLabels)  
val auROC = metrics.areaUnderROC()  
println("Area under ROC = " + auROC)
```


MLLib Example - Collaborative Filtering_

MLLib Example - Collaborative Filtering using ALS_

```
// Load and parse the data (userid,itemid,rating)
val data = sc.textFile("data/mllib/als/test.data")
val ratings = data.map(_.split(',').match
{
  case Array(user, item, rate) => Rating(user.toInt,
    item.toInt, rate.toDouble)
})

// Build the recommendation model using ALS
val rank = 10
val numIterations = 20
val model = ALS.train(ratings, rank, numIterations, 0.01)
```

MLLib Example - Collaborative Filtering using ALS

```
// Evaluate the model on rating data
val usersProducts = ratings.map {
  case Rating(user, product, rate) => (user, product) }

val predictions = model.predict(usersProducts).map {
  case Rating(user, product, rate) => ((user, product), rate)
}

val ratesAndPredictions = ratings.map {
  case Rating(user, product, rate) => ((user, product), rate)}
  .join(predictions)
val MSE = ratesAndPredictions.map {
  case ((user, product), (r1, r2)) => val err = (r1 - r2);
  err * err }.mean()
println("Mean Squared Error = " + MSE)
```


Use Cases

Data Loading

- In particular for huge amounts of external data
- Support for CSV, TSV, XML, JSON und other

```
case class User (id: java.util.UUID, name: String)

val users = sc.textFile("users.csv")
  .repartition(2*sc.defaultParallelism)
  .map(line => line.split(",") match { case Array(id,name) =>
    User(java.util.UUID.fromString(id), name)})

users.saveToCassandra("keyspace", "users")
```

Validation & Normalization

Validate consistency in a Cassandra database

- **syntactic**

- *Uniqueness (only relevant for columns not in the PK)*
- *Referential integrity*
- *Integrity of the duplicates*

- **semantic**

- *Business- or Application constraints*
- *e.g.: At least one genre per movies, a maximum of 10 tags per blog post*

Analyses (Joins, Transformations,..)

- Modelling, Mining, Transforming,
- Use Cases
 - *Recommendation*
 - *Fraud Detection*
 - *Link Analysis (Social Networks, Web)*
 - *Advertising*
 - *Data Stream Analytics (Spark Streaming)*
 - *Machine Learning (Spark ML)*

Schema Migration

- **Changes on existing tables**
 - *New table required when changing primary key*
 - *Otherwise changes could be performed in-place*
- **Creating new tables**
 - *data derived from existing tables*
 - *Support new queries*
- **Use the CassandraConnectors in Spark**

Questions?

Matthias Niehoff,

IT-Consultant

codecentric AG
Zeppelinstraße 2
76185 Karlsruhe, Germany

mobil: +49 (0) 172.1702676
matthias.niehoff@codecentric.de

www.codecentric.de
blog.codecentric.de

 matthiasniehoff

